

Part A

```
1 # 1.Import the data. The headlines will become your vectorized X matrix, and the labels indicate a binary classification (clickbait or not)
2
3 #mounting google drive files
4 from google.colab import drive
5 drive.mount('/content/drive')
```

Mounted at /content/drive

```
1 import pandas as pd
2 # Define file path
3 file_path = '/content/drive/My Drive/MLLAB4/text_training_data.csv'
4
5 # Load the CSV file into a DataFrame
6 data = pd.read_csv(file_path)
7 data.head()
```

	headline	label
0	MyBook Disk Drive Handles Lots of Easy Backups	not clickbait
1	CIT Posts Eighth Loss in a Row	not clickbait
2	Candy Carson Singing The "National Anthem" Is ...	clickbait
3	Why You Need To Stop What You're Doing And Dat...	clickbait
4	27 Times Adele Proved She's Actually The Reale...	clickbait

```
1 # 2.Convert the headline data into an X feature matrix using a simple bag of words approach
2
3 import numpy as np
4 from sklearn.feature_extraction.text import CountVectorizer
5
6 # Extract the headlines as a list of strings
7 docs = data['headline'].tolist()
8
9 count = CountVectorizer()
10 bag = count.fit_transform(docs)
11
12 # Check the vocabulary and transformed bag of words
13 print(count.vocabulary_)
14 print(count.get_feature_names_out())
15 print(bag)
16
```

```
{'mybook': 12118, 'disk': 5391, 'drive': 5708, 'handles': 8231, 'lots': 10820, 'of': 12660, 'easy': 5880, 'backups': 1668, 'cit': 3619,
['00' '000' '000th' ... 'ürümqi' 'şrı' 'şibenik']}
(0, 12118) 1
(0, 5391) 1
(0, 5708) 1
(0, 8231) 1
(0, 10820) 1
(0, 12660) 1
(0, 5880) 1
(0, 1668) 1
(1, 3619) 1
(1, 13925) 1
(1, 5965) 1
(1, 10815) 1
(1, 9123) 1
(1, 15606) 1
(2, 3049) 1
(2, 3163) 1
(2, 16581) 1
(2, 18225) 2
(2, 12214) 1
(2, 1116) 1
(2, 9547) 1
(2, 12744) 1
(2, 18258) 1
(2, 20213) 1
(2, 12264) 1
:
(24976, 19717) 1
```

```
(24976, 8583) 1
(24976, 19465) 1
(24976, 19681) 1
(24976, 4347) 1
(24976, 238) 1
(24977, 876) 1
(24977, 1433) 1
(24977, 7144) 1
(24977, 17846) 1
(24977, 13702) 1
(24977, 5749) 1
(24977, 18544) 1
(24977, 13810) 1
(24977, 18138) 1
(24977, 16975) 1
(24978, 9123) 1
(24978, 4836) 1
(24978, 9526) 1
(24978, 17281) 1
(24978, 14699) 1
(24978, 19083) 1
(24978, 1) 1
(24978, 11631) 1
(24978, 18458) 1
```

```
1 print(bag.toarray())
```

```
→ [[0 0 0 ... 0 0 0]
   [0 0 0 ... 0 0 0]
   [0 0 0 ... 0 0 0]
   ...
   [0 0 0 ... 0 0 0]
   [0 0 0 ... 0 0 0]
   [0 1 0 ... 0 0 0]]
```

```
1 # Logistical Regression using Bag of Words
2 import numpy as np
3 import pandas as pd
4 from sklearn.model_selection import train_test_split, GridSearchCV
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.metrics import classification_report, f1_score
7 from sklearn.preprocessing import LabelEncoder
8
9 # Use the existing bag-of-words matrix (X) and labels (y)
10 X = bag # Reuse the bag-of-words matrix created earlier
11 y = data['label'] # Assuming the label column contains string labels
12
13 # Encode the string labels into numeric values
14 label_encoder = LabelEncoder()
15 y_encoded = label_encoder.fit_transform(y) # 'clickbait' -> 1, 'not clickbait' -> 0
16
17 # Split data into training and test sets
18 X_train, X_test, y_train, y_test = train_test_split(X, y_encoded, test_size=0.2, random_state=42)
19
20 # Define logistic regression model
21 log_reg = LogisticRegression(max_iter=1000, random_state=42)
22
23 # Set up hyperparameter grid for C (inverse of regularization strength)
24 param_grid = {'C': [0.1, 1, 10, 100, 1000]}
25
26 # Use GridSearchCV to find the best hyperparameters
27 grid_search = GridSearchCV(log_reg, param_grid, scoring='f1', cv=5, n_jobs=-1)
28 grid_search.fit(X_train, y_train)
29
30 # Get the best model and parameters
31 best_model = grid_search.best_estimator_
32 best_params = grid_search.best_params_
33 print(f"Best parameters: {best_params}")
34
35 # Evaluate on the test set
36 y_pred = best_model.predict(X_test)
37 f1 = f1_score(y_test, y_pred)
38 print(f"F1 Score on the test set: {f1}")
39 print("\nClassification Report:")
40 print(classification_report(y_test, y_pred, target_names=label_encoder.classes_))
41
```

Best parameters: {'C': 10}
F1 Score on the test set: 0.9696391063586023

Classification Report:				
	precision	recall	f1-score	support
clickbait	0.97	0.96	0.97	2386
not clickbait	0.97	0.97	0.97	2610
accuracy			0.97	4996
macro avg	0.97	0.97	0.97	4996
weighted avg	0.97	0.97	0.97	4996

```

1 # N-grams
2 from sklearn.feature_extraction.text import CountVectorizer
3 from sklearn.pipeline import Pipeline
4 from sklearn.model_selection import GridSearchCV, train_test_split
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.metrics import classification_report, f1_score
7
8 # Split the dataset
9 X = data['headline'] # Original text data
10 y = label_encoder.transform(data['label']) # Numeric labels from previous steps
11
12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
13
14 # Create pipeline for logistic regression with n-grams
15 pipeline_ngram = Pipeline([
16     ('vectorizer', CountVectorizer(ngram_range=(1, 2))), # Use unigrams and bigrams
17     ('classifier', LogisticRegression(max_iter=1000, random_state=42))
18 ])
19
20 # Hyperparameter tuning
21 param_grid = {'classifier__C': [0.1, 1, 10, 100, 1000]} # Regularization strength
22 grid_search_ngram = GridSearchCV(pipeline_ngram, param_grid, scoring='f1', cv=5, n_jobs=-1)
23 grid_search_ngram.fit(X_train, y_train)
24
25 # Best model and evaluation
26 best_model_ngram = grid_search_ngram.best_estimator_
27 print(f"Best parameters for n-grams: {grid_search_ngram.best_params}")
28
29 y_pred_ngram = best_model_ngram.predict(X_test)
30 f1_ngram = f1_score(y_test, y_pred_ngram)
31 print(f"F1 Score for n-grams: {f1_ngram}")
32 print("\nClassification Report for n-grams:")
33 print(classification_report(y_test, y_pred_ngram))
34

```

Best parameters for n-grams: {'classifier__C': 100}
F1 Score for n-grams: 0.9688871922122543

Classification Report for n-grams:				
	precision	recall	f1-score	support
0	0.97	0.96	0.97	2386
1	0.97	0.97	0.97	2610
accuracy			0.97	4996
macro avg	0.97	0.97	0.97	4996
weighted avg	0.97	0.97	0.97	4996

```

1 # Stop-words
2 from sklearn.feature_extraction.text import CountVectorizer
3 from sklearn.pipeline import Pipeline
4 from sklearn.model_selection import GridSearchCV, train_test_split
5 from sklearn.linear_model import LogisticRegression
6 from sklearn.metrics import classification_report, f1_score
7
8 # Split the dataset
9 X = data['headline'] # Original text data
10 y = label_encoder.transform(data['label']) # Numeric labels from previous steps
11
12 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
13
14 # Create pipeline for logistic regression with stop-words removal
15 pipeline_stopwords = Pipeline([

```

```

16 ('vectorizer', CountVectorizer(stop_words='english')), # Removes common stop words
17 ('classifier', LogisticRegression(max_iter=1000, random_state=42))
18 ])
19
20 # Hyperparameter tuning
21 param_grid = {'classifier__C': [0.1, 1, 10, 100, 1000]} # Regularization strength
22 grid_search_stopwords = GridSearchCV(pipeline_stopwords, param_grid, scoring='f1', cv=5, n_jobs=-1)
23 grid_search_stopwords.fit(X_train, y_train)
24
25 # Best model and evaluation
26 best_model_stopwords = grid_search_stopwords.best_estimator_
27 print(f"Best parameters for stop-words: {grid_search_stopwords.best_params_}")
28
29 y_pred_stopwords = best_model_stopwords.predict(X_test)
30 f1_stopwords = f1_score(y_test, y_pred_stopwords)
31 print(f"F1 Score for stop-words method: {f1_stopwords}")
32 print("\nClassification Report for stop-words:")
33 print(classification_report(y_test, y_pred_stopwords))
34

```

```

→ Best parameters for stop-words: {'classifier__C': 1}
F1 Score for stop-words method: 0.9512937595129376

```

```

Classification Report for stop-words:

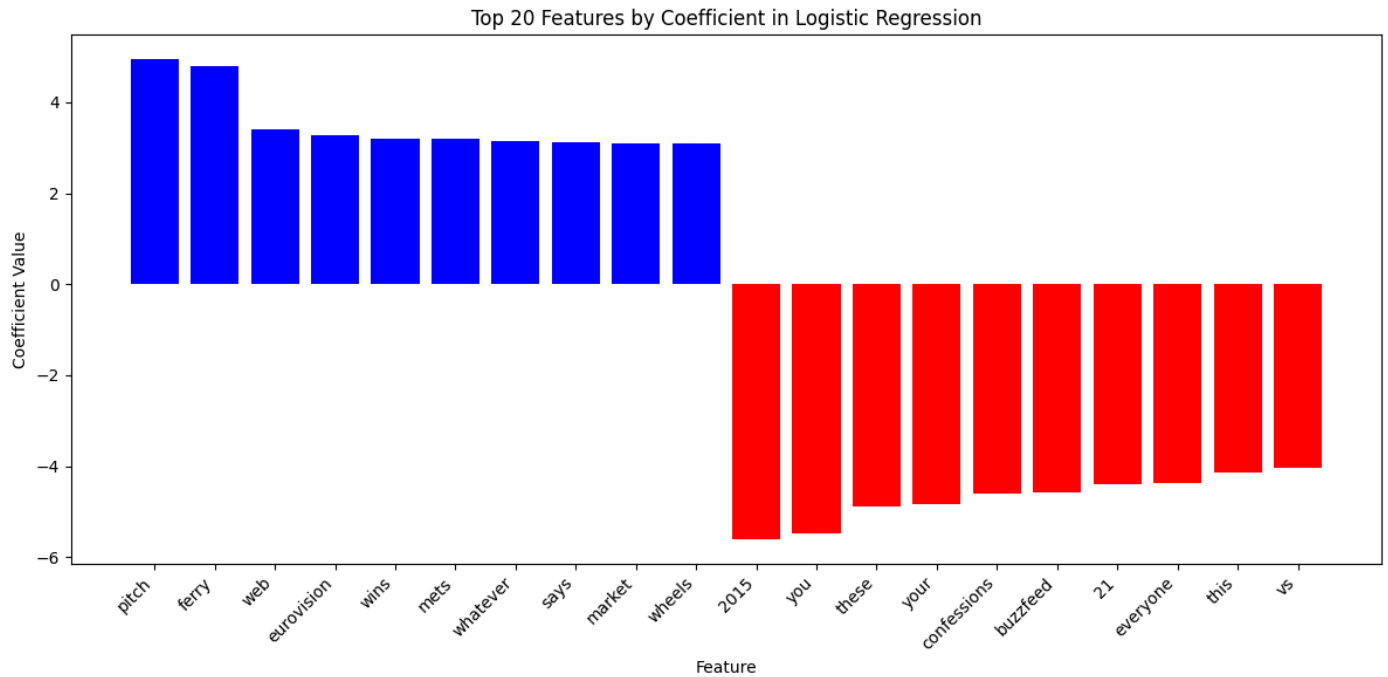
```

	precision	recall	f1-score	support
0	0.95	0.94	0.95	2386
1	0.94	0.96	0.95	2610
accuracy			0.95	4996
macro avg	0.95	0.95	0.95	4996
weighted avg	0.95	0.95	0.95	4996

```

1 #bag of words coefficients
2 import matplotlib.pyplot as plt
3 import pandas as pd
4
5 # Get the feature names from the CountVectorizer
6 feature_names = count.get_feature_names_out() # Use the 'count' object from earlier
7
8 # Get the coefficients from the best model
9 coefficients = best_model.coef_[0] # Coefficients for the logistic regression model
10
11 # Create a DataFrame to map coefficients to feature names
12 coef_df = pd.DataFrame({'Feature': feature_names, 'Coefficient': coefficients})
13
14 # Split into positive and negative coefficients
15 positive_features = coef_df[coef_df['Coefficient'] > 0].sort_values(by='Coefficient', ascending=False).head(10)
16 negative_features = coef_df[coef_df['Coefficient'] < 0].sort_values(by='Coefficient').head(10)
17
18 # Combine the top positive and negative coefficients
19 top_features = pd.concat([positive_features, negative_features])
20
21 # Plot the top features
22 plt.figure(figsize=(12, 6))
23
24 # Assign colors: red for negative coefficients, blue for positive
25 colors = ['blue' if coef > 0 else 'red' for coef in top_features['Coefficient']]
26
27 # Plot vertical bars
28 plt.bar(top_features['Feature'], top_features['Coefficient'], color=colors)
29
30 # Add labels and title
31 plt.ylabel('Coefficient Value')
32 plt.xlabel('Feature')
33 plt.title('Top 20 Features by Coefficient in Logistic Regression')
34 plt.xticks(rotation=45, ha='right') # Rotate x-axis labels for better readability
35 plt.tight_layout() # Adjust layout to prevent overlap
36 plt.show()
37

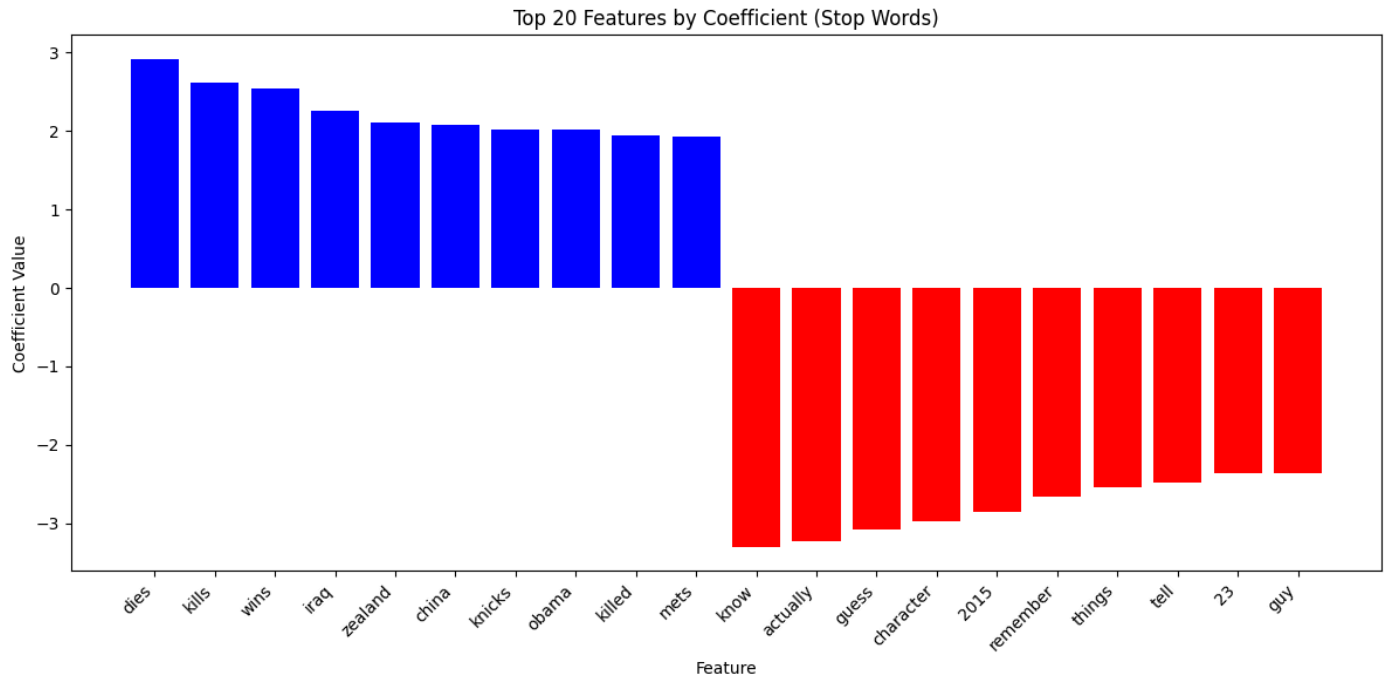
```



```

1 #stop-words coefficients
2 # Stop words vectorization
3 from sklearn.feature_extraction.text import CountVectorizer
4
5 # Vectorization with stop words
6 stop_words_vectorizer = CountVectorizer(stop_words='english') # Remove common stop words
7 bag_stop_words = stop_words_vectorizer.fit_transform(docs)
8
9 # Train logistic regression model
10 log_reg_stop_words = LogisticRegression(max_iter=1000, random_state=42)
11 log_reg_stop_words.fit(bag_stop_words, y_encoded)
12
13 # Get feature names and coefficients
14 feature_names_stop_words = stop_words_vectorizer.get_feature_names_out()
15 coefficients_stop_words = log_reg_stop_words.coef_[0]
16
17 # Create DataFrame for coefficients
18 coef_df_stop_words = pd.DataFrame({
19     'Feature': feature_names_stop_words,
20     'Coefficient': coefficients_stop_words
21 })
22
23 # Sort and select top 20 features
24 positive_features_stop = coef_df_stop_words[coef_df_stop_words['Coefficient'] > 0]\
25     .sort_values(by='Coefficient', ascending=False).head(10)
26 negative_features_stop = coef_df_stop_words[coef_df_stop_words['Coefficient'] < 0]\
27     .sort_values(by='Coefficient').head(10)
28
29 top_features_stop_words = pd.concat([positive_features_stop, negative_features_stop])
30
31 # Plot top 20 coefficients
32 plt.figure(figsize=(12, 6))
33 colors = ['blue' if coef > 0 else 'red' for coef in top_features_stop_words['Coefficient']]
34 plt.bar(top_features_stop_words['Feature'], top_features_stop_words['Coefficient'], color=colors)
35 plt.ylabel('Coefficient Value')
36 plt.xlabel('Feature')
37 plt.title('Top 20 Features by Coefficient (Stop Words)')
38 plt.xticks(rotation=45, ha='right')
39 plt.tight_layout()
40 plt.show()
41

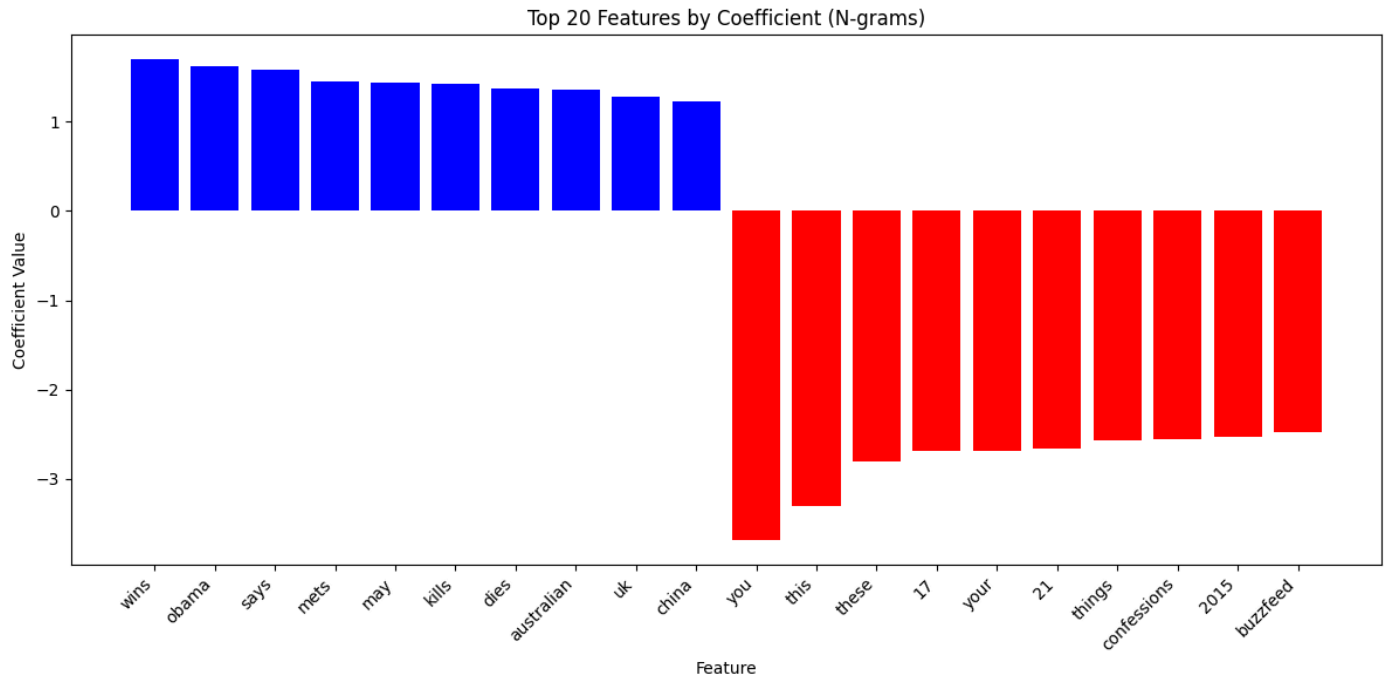
```



```

1 #n-grams coefficients
2 # N-grams vectorization
3 ngram_vectorizer = CountVectorizer(ngram_range=(1, 2)) # Unigrams and bigrams
4 bag_ngrams = ngram_vectorizer.fit_transform(docs)
5
6 # Train logistic regression model
7 log_reg_ngrams = LogisticRegression(max_iter=1000, random_state=42)
8 log_reg_ngrams.fit(bag_ngrams, y_encoded)
9
10 # Get feature names and coefficients
11 feature_names_ngrams = ngram_vectorizer.get_feature_names_out()
12 coefficients_ngrams = log_reg_ngrams.coef_[0]
13
14 # Create DataFrame for coefficients
15 coef_df_ngrams = pd.DataFrame({
16     'Feature': feature_names_ngrams,
17     'Coefficient': coefficients_ngrams
18 })
19
20 # Sort and select top 20 features
21 positive_features_ngrams = coef_df_ngrams[coef_df_ngrams['Coefficient'] > 0]\
22     .sort_values(by='Coefficient', ascending=False).head(10)
23 negative_features_ngrams = coef_df_ngrams[coef_df_ngrams['Coefficient'] < 0]\
24     .sort_values(by='Coefficient').head(10)
25
26 top_features_ngrams = pd.concat([positive_features_ngrams, negative_features_ngrams])
27
28 # Plot top 20 coefficients
29 plt.figure(figsize=(12, 6))
30 colors = ['blue' if coef > 0 else 'red' for coef in top_features_ngrams['Coefficient']]
31 plt.bar(top_features_ngrams['Feature'], top_features_ngrams['Coefficient'], color=colors)
32 plt.ylabel('Coefficient Value')
33 plt.xlabel('Feature')
34 plt.title('Top 20 Features by Coefficient (N-grams)')
35 plt.xticks(rotation=45, ha='right')
36 plt.tight_layout()
37 plt.show()
38

```



Of the three models, the bag of words model appears to perform the best in terms of f-1 score, while the stop-words model performs the worst. Precision and recall for the bag of words model and the n-grams model are identical, but the f-1 score for the n-grams model is slightly worse. In terms of coefficients, it seems like all three models generally attribute the same coefficients to each feature, with the only exception being the stop-words model, which doesn't attribute much importance to the 'you' feature.

Part B

```
1 url = 'http://vincentarelbundock.github.io/Rdatasets/csv/datasets/iris.csv'
2 data = pd.read_csv(url)
3 data.head()
```



	rownames	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
0	1	5.1	3.5	1.4	0.2	setosa
1	2	4.9	3.0	1.4	0.2	setosa
2	3	4.7	3.2	1.3	0.2	setosa
3	4	4.6	3.1	1.5	0.2	setosa
4	5	5.0	3.6	1.4	0.2	setosa

```
1 import pandas as pd
2
3 # Example: Load your dataset (replace with actual file path or source)
4 data = pd.read_csv(url)
5
6 # Number of features (assuming the target column is the last one)
7 n_features = data.shape[1] - 1
8 print(f"Number of features: {n_features}")
9
```



Number of features: 5

```
1 # Load libraries
2 import numpy as np
3 !pip install tensorflow
4 import pandas as pd
5 import tensorflow.keras as keras
```

```

6 from tensorflow.keras.models import Sequential
7 from tensorflow.keras.layers import Dense, Dropout, Activation
8 from tensorflow.keras.optimizers import SGD
9
10
11 # The core data structure of Keras is a model, a way to organize layers
12 # Here we are using a Sequential model architecture
13 model = Sequential()

```

Requirement already satisfied: tensorflow in /usr/local/lib/python3.10/dist-packages (2.17.1)

Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.4.0)

Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.6.3)

Requirement already satisfied: flatbuffers>=24.3.25 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (24.3.25)

Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.6.0)

Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.2.0)

Requirement already satisfied: h5py>=3.10.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.12.1)

Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (18.1.1)

Requirement already satisfied: ml-dtypes<0.5.0,>=0.3.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.4.1)

Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.4.0)

Requirement already satisfied: packaging in /usr/local/lib/python3.10/dist-packages (from tensorflow) (24.2)

Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<5.0.0dev,>=3.20.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (4.21.0)

Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.32.3)

Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from tensorflow) (75.1.0)

Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)

Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.5.0)

Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (4.12.2)

Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.16.0)

Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.68.0)

Requirement already satisfied: tensorboard<2.18,>=2.17 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (2.17.1)

Requirement already satisfied: keras>=3.2.0 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (3.5.0)

Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (0.37.1)

Requirement already satisfied: numpy<2.0.0,>=1.23.5 in /usr/local/lib/python3.10/dist-packages (from tensorflow) (1.26.4)

Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.10/dist-packages (from astunparse>=1.6.0->tensorflow) (0.45.0)

Requirement already satisfied: rich in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (13.9.4)

Requirement already satisfied: namex in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (0.0.8)

Requirement already satisfied: optree in /usr/local/lib/python3.10/dist-packages (from keras>=3.2.0->tensorflow) (0.13.1)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.4.0)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (3.10)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (2.2.3)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorflow) (2025.1.1)

Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17->tensorflow) (3.7)

Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17->tensorflow) (0.17.0)

Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from tensorboard<2.18,>=2.17->tensorflow) (3.0.6)

Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.10/dist-packages (from werkzeug>=1.0.1->tensorboard<2.18,>=2.17->tensorflow) (3.0.2)

Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0->tensorflow) (3.0.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.10/dist-packages (from rich->keras>=3.2.0->tensorflow) (2.19.2)

Requirement already satisfied: mdurl~0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-it-py>=2.2.0->rich->keras>=3.2.0->tensorflow) (0.1.2)

```

1
2 #NN 2 layers 16 neurons
3 from sklearn.model_selection import train_test_split
4 from tensorflow.keras.utils import to_categorical
5 from tensorflow.keras.models import Sequential
6 from tensorflow.keras.layers import Dense
7 from tensorflow.keras.optimizers import SGD
8
9 # Assume 'data' is your pandas DataFrame with features and target
10 # Split features (X) and target (y)
11 X = data.drop(columns=['Species']).values # Features
12 y = data['Species'].values # Target
13
14 # Convert string labels in y to one-hot encoding directly
15 unique_classes = list(set(y)) # Extract unique class labels
16 num_classes = len(unique_classes) # Number of unique classes
17 class_to_index = {cls: i for i, cls in enumerate(unique_classes)} # Map classes to integers
18 y = [class_to_index[label] for label in y] # Convert labels to integers
19
20 # One-hot encode the y data using to_categorical()
21 y = to_categorical(y, num_classes=num_classes)
22
23 # Split into train and test sets
24 x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
25
26 # Build the model
27 model = Sequential()
28 model.add(Dense(16, activation='relu', input_dim=X.shape[1])) # First hidden layer
29 model.add(Dense(16, activation='relu')) # Second hidden layer
30 model.add(Dense(num_classes, activation='softmax')) # Output layer
31

```



```

32 # Optimize using SGD with a learning rate
33 sgd = SGD(learning_rate=0.01)
34 model.compile(loss='categorical_crossentropy',
35               optimizer=sgd,
36               metrics=['accuracy'])
37
38 # Train the model
39 model.fit(x_train, y_train,
40           epochs=20,
41           batch_size=128)
42
43 # Evaluate the model
44 score = model.evaluate(x_test, y_test, batch_size=128) # Extract loss and accuracy
45 print(f"Loss: {score[0]}, Accuracy: {score[1]}")
46

```

⚠ /usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to the `__init__` method of `Dense` layer. You should pass it to the `add` method of the `Model` class instead.

```

Epoch 1/20
1/1 ----- 3s 3s/step - accuracy: 0.3250 - loss: 6.8055
Epoch 2/20
1/1 ----- 0s 149ms/step - accuracy: 0.3333 - loss: 5.2705
Epoch 3/20
1/1 ----- 0s 315ms/step - accuracy: 0.3250 - loss: 12.2882
Epoch 4/20
1/1 ----- 0s 267ms/step - accuracy: 0.3417 - loss: 7.6394
Epoch 5/20
1/1 ----- 0s 118ms/step - accuracy: 0.3250 - loss: 6.2564
Epoch 6/20
1/1 ----- 0s 130ms/step - accuracy: 0.3333 - loss: 1.4733
Epoch 7/20
1/1 ----- 0s 123ms/step - accuracy: 0.3500 - loss: 1.6632
Epoch 8/20
1/1 ----- 0s 209ms/step - accuracy: 0.3333 - loss: 1.9135
Epoch 9/20
1/1 ----- 0s 331ms/step - accuracy: 0.3417 - loss: 1.0598
Epoch 10/20
1/1 ----- 0s 287ms/step - accuracy: 0.3583 - loss: 1.1099
Epoch 11/20
1/1 ----- 0s 311ms/step - accuracy: 0.3500 - loss: 1.2615
Epoch 12/20
1/1 ----- 0s 267ms/step - accuracy: 0.3667 - loss: 1.1979
Epoch 13/20
1/1 ----- 0s 154ms/step - accuracy: 0.3667 - loss: 1.2466
Epoch 14/20
1/1 ----- 0s 146ms/step - accuracy: 0.3417 - loss: 1.0470
Epoch 15/20
1/1 ----- 0s 155ms/step - accuracy: 0.3750 - loss: 1.0507
Epoch 16/20
1/1 ----- 0s 162ms/step - accuracy: 0.3833 - loss: 1.0324
Epoch 17/20
1/1 ----- 0s 153ms/step - accuracy: 0.3750 - loss: 1.0725
Epoch 18/20
1/1 ----- 0s 288ms/step - accuracy: 0.3917 - loss: 1.0589
Epoch 19/20
1/1 ----- 0s 100ms/step - accuracy: 0.3750 - loss: 1.0805
Epoch 20/20
1/1 ----- 0s 98ms/step - accuracy: 0.5833 - loss: 0.9391
1/1 ----- 1s 725ms/step - accuracy: 0.3667 - loss: 1.2851
Loss: 1.2851481437683105, Accuracy: 0.36666667461395264

```

```

1 #NN 4 layers 32 neurons
2 from sklearn.model_selection import train_test_split
3 from tensorflow.keras.utils import to_categorical
4 from tensorflow.keras.models import Sequential
5 from tensorflow.keras.layers import Dense
6 from tensorflow.keras.optimizers import SGD
7
8 # Assume 'data' is your pandas DataFrame with features and target
9 # Split features (X) and target (y)
10 X = data.drop(columns=['Species']).values # Features
11 y = data['Species'].values # Target
12
13 # Convert string labels in y to one-hot encoding directly
14 unique_classes = list(set(y)) # Extract unique class labels
15 num_classes = len(unique_classes) # Number of unique classes
16 class_to_index = {cls: i for i, cls in enumerate(unique_classes)} # Map classes to integers
17 y = [class_to_index[label] for label in y] # Convert labels to integers
18

```

```

19 # One-hot encode the y data using to_categorical()
20 y = to_categorical(y, num_classes=num_classes)
21
22 # Split into train and test sets
23 x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
24
25 # Build the model
26 model = Sequential()
27 model.add(Dense(32, activation='relu', input_dim=X.shape[1])) # First hidden layer
28 model.add(Dense(32, activation='relu')) # Second hidden layer
29 model.add(Dense(32, activation='relu')) # Third hidden layer
30 model.add(Dense(32, activation='relu')) # Fourth hidden layer
31 model.add(Dense(num_classes, activation='softmax')) # Output layer
32
33 # Optimize using SGD with a learning rate
34 sgd = SGD(learning_rate=0.01)
35 model.compile(loss='categorical_crossentropy',
36               optimizer=sgd,
37               metrics=['accuracy'])
38
39 # Train the model
40 model.fit(x_train, y_train,
41          epochs=20,
42          batch_size=128)
43
44 # Evaluate the model
45 score = model.evaluate(x_test, y_test, batch_size=128) # Extract loss and accuracy
46 print(f"Loss: {score[0]}, Accuracy: {score[1]}")
47

```

```

Epoch 1/20
/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
1/1 ————— 1s 623ms/step - accuracy: 0.3417 - loss: 5.4965
Epoch 2/20
1/1 ————— 0s 54ms/step - accuracy: 0.3417 - loss: 1.8527
Epoch 3/20
1/1 ————— 0s 32ms/step - accuracy: 0.3417 - loss: 1.2210
Epoch 4/20
1/1 ————— 0s 32ms/step - accuracy: 0.3250 - loss: 1.1889
Epoch 5/20
1/1 ————— 0s 57ms/step - accuracy: 0.4750 - loss: 0.9970
Epoch 6/20
1/1 ————— 0s 32ms/step - accuracy: 0.3417 - loss: 1.0022
Epoch 7/20
1/1 ————— 0s 34ms/step - accuracy: 0.3500 - loss: 1.0031
Epoch 8/20
1/1 ————— 0s 34ms/step - accuracy: 0.3583 - loss: 1.0149
Epoch 9/20
1/1 ————— 0s 30ms/step - accuracy: 0.3417 - loss: 0.9949
Epoch 10/20
1/1 ————— 0s 29ms/step - accuracy: 0.3750 - loss: 1.0073
Epoch 11/20
1/1 ————— 0s 32ms/step - accuracy: 0.3750 - loss: 0.9776
Epoch 12/20
1/1 ————— 0s 39ms/step - accuracy: 0.3750 - loss: 0.9838
Epoch 13/20
1/1 ————— 0s 36ms/step - accuracy: 0.4333 - loss: 0.9539
Epoch 14/20
1/1 ————— 0s 56ms/step - accuracy: 0.4000 - loss: 0.9517
Epoch 15/20
1/1 ————— 0s 34ms/step - accuracy: 0.5833 - loss: 0.9336
Epoch 16/20
1/1 ————— 0s 46ms/step - accuracy: 0.4417 - loss: 0.9278
Epoch 17/20
1/1 ————— 0s 41ms/step - accuracy: 0.5917 - loss: 0.9143
Epoch 18/20
1/1 ————— 0s 59ms/step - accuracy: 0.4583 - loss: 0.9060
Epoch 19/20
1/1 ————— 0s 41ms/step - accuracy: 0.6250 - loss: 0.8956
Epoch 20/20
1/1 ————— 0s 58ms/step - accuracy: 0.4750 - loss: 0.8907
WARNING:tensorflow:5 out of the last 5 calls to <function TensorFlowTrainer.make_test_function.<locals>.one_step_on_iterator at 0x7e1f4b
1/1 ————— 0s 215ms/step - accuracy: 0.6667 - loss: 0.8665
Loss: 0.866498589515686, Accuracy: 0.6666666665348816

```

```

1 #NN 8 layers 64 neurons
2 from sklearn.model_selection import train_test_split
3 from tensorflow.keras.utils import to_categorical

```

```

4 from tensorflow.keras.models import Sequential
5 from tensorflow.keras.layers import Dense
6 from tensorflow.keras.optimizers import SGD
7
8 # Assume 'data' is your pandas DataFrame with features and target
9 # Split features (X) and target (y)
10 X = data.drop(columns=['Species']).values # Features
11 y = data['Species'].values # Target
12
13 # Convert string labels in y to one-hot encoding directly
14 unique_classes = list(set(y)) # Extract unique class labels
15 num_classes = len(unique_classes) # Number of unique classes
16 class_to_index = {cls: i for i, cls in enumerate(unique_classes)} # Map classes to integers
17 y = [class_to_index[label] for label in y] # Convert labels to integers
18
19 # One-hot encode the y data using to_categorical()
20 y = to_categorical(y, num_classes=num_classes)
21
22 # Split into train and test sets
23 x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
24
25 # Build the model
26 model = Sequential()
27 model.add(Dense(64, activation='relu', input_dim=X.shape[1])) # First hidden layer
28 model.add(Dense(64, activation='relu')) # Second hidden layer
29 model.add(Dense(64, activation='relu')) # Third hidden layer
30 model.add(Dense(64, activation='relu')) # Fourth hidden layer
31 model.add(Dense(64, activation='relu')) # Fifth hidden layer
32 model.add(Dense(64, activation='relu')) # Sixth hidden layer
33 model.add(Dense(64, activation='relu')) # Seventh hidden layer
34 model.add(Dense(64, activation='relu')) # Eighth hidden layer
35 model.add(Dense(num_classes, activation='softmax')) # Output layer
36
37 # Optimize using SGD with a learning rate
38 sgd = SGD(learning_rate=0.01)
39 model.compile(loss='categorical_crossentropy',
40               optimizer=sgd,
41               metrics=['accuracy'])
42
43 # Train the model
44 model.fit(x_train, y_train,
45          epochs=20,
46          batch_size=128)
47
48 # Evaluate the model
49 score = model.evaluate(x_test, y_test, batch_size=128) # Extract loss and accuracy
50 print(f"Loss: {score[0]}, Accuracy: {score[1]}")
51

```

```

Epoch 1/20
/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
1/1 ━━━━━━━━━━━ 1s 1s/step - accuracy: 0.3250 - loss: 1.2194
Epoch 2/20
1/1 ━━━━━━━━━━━ 0s 51ms/step - accuracy: 0.3250 - loss: 0.9934
Epoch 3/20
1/1 ━━━━━━━━━━━ 0s 56ms/step - accuracy: 0.3250 - loss: 0.9716
Epoch 4/20
1/1 ━━━━━━━━━━━ 0s 40ms/step - accuracy: 0.3250 - loss: 0.9664
Epoch 5/20
1/1 ━━━━━━━━━━━ 0s 45ms/step - accuracy: 0.3250 - loss: 0.9630
Epoch 6/20
1/1 ━━━━━━━━━━━ 0s 56ms/step - accuracy: 0.3250 - loss: 0.9620
Epoch 7/20
1/1 ━━━━━━━━━━━ 0s 44ms/step - accuracy: 0.3250 - loss: 0.9567
Epoch 8/20
1/1 ━━━━━━━━━━━ 0s 56ms/step - accuracy: 0.3417 - loss: 0.9539
Epoch 9/20
1/1 ━━━━━━━━━━━ 0s 51ms/step - accuracy: 0.3750 - loss: 0.9512
Epoch 10/20
1/1 ━━━━━━━━━━━ 0s 59ms/step - accuracy: 0.3917 - loss: 0.9476
Epoch 11/20
1/1 ━━━━━━━━━━━ 0s 131ms/step - accuracy: 0.4000 - loss: 0.9446
Epoch 12/20
1/1 ━━━━━━━━━━━ 0s 70ms/step - accuracy: 0.4000 - loss: 0.9428
Epoch 13/20
1/1 ━━━━━━━━━━━ 0s 47ms/step - accuracy: 0.4083 - loss: 0.9374
Epoch 14/20
1/1 ━━━━━━━━━━━ 0s 59ms/step - accuracy: 0.4167 - loss: 0.9329
Epoch 15/20

```

```
1/1 ————— 0s 60ms/step - accuracy: 0.4167 - loss: 0.9296
Epoch 16/20
1/1 ————— 0s 58ms/step - accuracy: 0.4167 - loss: 0.9273
Epoch 17/20
1/1 ————— 0s 58ms/step - accuracy: 0.4167 - loss: 0.9292
Epoch 18/20
1/1 ————— 0s 69ms/step - accuracy: 0.4417 - loss: 0.9225
Epoch 19/20
1/1 ————— 0s 58ms/step - accuracy: 0.4250 - loss: 0.9224
Epoch 20/20
1/1 ————— 0s 76ms/step - accuracy: 0.4500 - loss: 0.9126
1/1 ————— 0s 295ms/step - accuracy: 0.4667 - loss: 0.8812
Loss: 0.8812350630760193, Accuracy: 0.46666666865348816
```

Overall, it appears that adding neurons and layers does in fact improve the accuracy of the model, with the second and third models observing higher accuracy scores than the first model (less layers and neurons). However, one thing I did notice is that the accuracy scores can fluctuate between the second and third models, which indicates that for small datasets there may be a local minima/maxima where adding a certain number of layers or neurons becomes futile in improving accuracy.